

Evaluating Robustness and Generalization in LLMs under Adversarial and Real-World Conditions

Yigit Demirsan

Department of Artificial Intelligence
Ozyegin University
Istanbul, Turkey
yigit.demirsan@ozu.edu.tr

Olcay Taner Yildiz

Department of Computer Science and Engineering
Ozyegin University
Istanbul, Turkey
olcay.yildiz@ozyegin.edu.tr

Abstract—In this paper, we present a comprehensive survey on the robustness and generalization of large language models (LLMs) under adversarial and real-world conditions, covering research published between 2020 and 2025. The survey systematically analyzes adversarial vulnerabilities, including prompt injection attacks and word level perturbations, and examines the extent to which LLMs can generalize effectively across diverse domains and unseen tasks. Evaluation frameworks such as CheckList, TextAttack, and Robustness Gym are thoroughly reviewed, along with adversarial robustness benchmarks including ANLI, AdvGLUE, and datasets produced through the Dynabench platform. Additionally, key metrics used to measure model robustness and generalization—such as adversarial accuracy, robustness degradation, and performance consistency are critically evaluated.

The robustness characteristics of state-of-the-art models, including GPT-4, Claude, and LLaMA, are compared and synthesized, highlighting their strengths and vulnerabilities in adversarial scenarios and under distribution shifts. Despite significant advancements, the findings indicate persistent vulnerabilities across all models under targeted adversarial attacks or domain variation. Furthermore, the paper reviews and assesses robustness improvement techniques, including adversarial training, data augmentation strategies, instruction tuning, and constitutional AI, identifying their effectiveness as well as limitations. Ultimately, this survey emphasizes the necessity for comprehensive robustness evaluations that reflect diverse adversarial and practical stressors, underscoring the critical role of robust language models for safe and reliable real-world deployments.

Index Terms—Large Language Models, Natural Language Processing, Transformers, Deep Learning

I. INTRODUCTION

Large Language Models (LLMs) have achieved impressive performance on standard benchmarks, but their robustness and generalization remain critical concerns – especially when faced with adversarial inputs or shifting real-world conditions. Traditional evaluation by held-out accuracy often overestimates true performance. Robustness in this context refers to an LLM’s resilience to perturbed or malicious inputs (e.g. adversarial examples or prompt attacks) without degrading its output quality, while generalization denotes consistent performance on tasks, domains, or distributions not seen in training. Recent research (2020–2025) has revealed that even state-of-the-art LLMs are brittle in surprising ways: for example, swapping a single word with a synonym can completely flip a model’s prediction despite unchanged meaning. Likewise, cleverly crafted

“jailbreak” prompts can trick even aligned conversational models into producing disallowed content. These failure modes underscore the need for comprehensive evaluation frameworks beyond standard accuracy. In this survey, we review adversarial robustness and generalization evaluation for LLMs, covering recent high-quality papers, evaluation frameworks (e.g. CheckList, TextAttack), robustness benchmarks, metrics, findings on state-of-the-art models (GPT-4, LLaMA, Claude, Gemini, etc.), and techniques to improve robustness.

II. BACKGROUND AND DEFINITIONS

Recent advances in natural language processing (NLP) have led to the development of large language models (LLMs) demonstrating remarkable capabilities across numerous linguistic tasks. These models, typically built on Transformer architectures such as GPT, BERT, and T5, are pretrained on extensive datasets and exhibit impressive zero-shot and few-shot performance. However, empirical evidence increasingly indicates potential vulnerabilities and generalization failures when these models encounter conditions deviating from their training distributions.

Robustness, within the context of LLMs, refers to a model’s capacity to sustain consistent and reliable performance under perturbed, noisy, adversarial, or intentionally deceptive inputs. A robust model should correctly handle semantically equivalent inputs despite superficial alterations, including synonym substitutions, typographical errors, or syntactic variations. Critically, robustness also entails resistance to adversarial inputs specifically crafted to provoke incorrect or harmful responses, such as “jailbreak” prompts or misleading instructions, thereby posing significant risks for real-world applications.

Generalization, on the other hand, denotes a model’s ability to transfer learned knowledge effectively to novel tasks, domains, or distributions without explicit retraining. This encompasses handling domain shifts, temporal shifts (e.g., outdated information), and user inputs distinct from training and benchmark datasets. Effective generalization is essential for reliable long-term deployment and trustworthiness in diverse practical scenarios.

Adversarial attacks in NLP represent diverse input manipulations devised explicitly to deceive or mislead models.

These attacks can manifest as character-level perturbations, paraphrased questions, semantic modifications, or sophisticated instruction-based exploits compromising model safety. Compared to adversarial scenarios in computer vision, attacks in NLP are discrete and often semantically nuanced, complicating their detection and mitigation.

To address these robustness and generalization challenges, the NLP community has developed specialized **evaluation frameworks, benchmarks, and metrics**. Tools such as *Check-List*, *TextAttack*, and *Robustness Gym* offer systematic approaches to behavioral testing, while benchmarks like *ANLI*, *AdvGLUE*, and *HELM* provide targeted datasets aimed at revealing model weaknesses.

This survey builds upon these foundational concepts, systematically reviewing the literature to synthesize current methodologies, findings, and proposed strategies for evaluating and enhancing LLM robustness and generalization. The survey identifies key insights, examines the efficacy of various approaches, and outlines future research directions toward more robust and trustworthy language models.

III. ADVERSARIAL VULNERABILITIES OF LLMs

Large Language Models have achieved remarkable accuracy on many NLP tasks, but a growing body of research reveals that their performance can degrade sharply under adversarial input perturbations. Recent experiments highlight how synonym substitutions, character-level noise (typos), and paraphrasing can significantly impair LLM outputs, and how model scale and architecture affect robustness. Moreover, specially crafted prompt injection (“jailbreak”) attacks can induce undesired behaviors even in state-of-the-art aligned models. Below, we survey representative findings (2020–2025) with figures illustrating these vulnerabilities.

A. Word-Level Attacks: Synonyms and Paraphrasing

Attack success rates for synonym substitution attacks on five text classification tasks, plotted against model size (T5 encoder-decoder models from 80M up to 11B parameters). Higher Attack Success Rate (ASR) means the adversary was more often able to flip the model’s prediction by replacing input words with synonyms. Larger T5 models tend to have lower ASR (better robustness) on most tasks, but some tasks (e.g. MRPC in blue, SST-2 in red) remain highly vulnerable even at 3B–11B scale. Smaller models show near-maximum ASR, indicating they can be easily fooled by slight wording changes. Word-level perturbations like synonym replacements can dramatically reduce an LLM’s accuracy without changing the input’s meaning. For example, Yang et al. (2024) find that simply swapping important words with synonyms causes high error rates across multiple datasets. The 80M-parameter T5 model is almost entirely compromised (ASR 0.9) on several tasks, whereas the 11B model resists more (ASR drops to 0.5 on SST-2, 0.1 on AGNews). This trend suggests larger models grasp language nuances better and are harder to trick with synonyms, though not invulnerable. Interestingly, model architecture also matters: decoder-only LLMs (OPT,

LLaMA) were found to have lower ASR than an encoder-decoder (T5) at similar scales, indicating higher robustness under these attacks. Beyond one-for-one word swaps, even more subtle paraphrasing can degrade LLM performance. Recent work by Arora et al. (2025) shows that rephrasing inputs in different sociodemographic styles (e.g. formality or dialect) significantly impacts accuracy, revealing brittleness to linguistic variation. In other cases, adding innocuous sentences or using an alternative phrasing has led QA systems to answer incorrectly (a phenomenon noted as early as Jia Liang 2017). These findings underscore that LLMs, despite their fluency, often latch onto specific wording – when that wording changes, the models can become confused or misled.




Fig. 1. Attack Success Rates for various datasets

B. Character-Level Attacks: Typos and Noise

Minor character-level perturbations – such as random typos, letter swaps, or deletions – can have an outsized effect on LLM outputs. Unlike humans, who easily read through minor spelling errors, models may misinterpret or lose context entirely. For instance, a recent robustness evaluation of GPT series models showed substantial accuracy drops on sentiment analysis when inputs contained random deletions of characters. GPT-4’s accuracy on IMDB reviews fell from 94–96% to 71–79% under a character deletion attack, a loss of over 20 percentage points. Other GPT-3.5/4 variants and open-source models saw similarly large degradations. All models tested “show a notable decrease in accuracy when exposed to character-level perturbations”, confirming that even small typos can derail LLM predictions. One intriguing observation is that the most advanced model (GPT-4) had the highest baseline accuracy yet also suffered one of the largest absolute drops under character attacks. In contrast, GPT-3.5-turbo, which started with slightly lower clean accuracy, dropped a bit less in some cases. This hints at a possible trade-off: models that are highly optimized for accuracy might be more sensitive to distribution shifts like misspellings. Prior studies on BERT and other classifiers noted similar vulnerability to typos (e.g. HotFlip attacks), often causing the model’s

confidence to collapse on perturbed inputs. Overall, character-level robustness remains a key weakness for LLMs – an adversary could exploit this by introducing imperceptible typos that dramatically change the model’s response.

C. Influence of Model Size on Robustness

Does making models larger help them resist adversarial attacks? The evidence is mixed. On one hand, scaling up generally improves a model’s linguistic understanding and might reduce susceptibility to simple attacks. For example, as shown earlier, a 3B or 11B parameter T5 endured synonym perturbations better than a 60M version. Likewise, OpenAI’s GPT-4 (with tens of billions of parameters) outperforms GPT-3.5 in many robust reasoning benchmarks, and it also detected significantly more adversarial/jailbreak prompts than its predecessor (discussed below). These cases suggest that larger models develop more “global” language representations, making them a bit harder to fool with local edits. However, larger size alone is not a panacea. Some studies have found that even the biggest LLMs still fail on cleverly constructed inputs. For instance, instruction-tuning (an extra training step for LLMs) can paradoxically make a model more vulnerable to word-level attacks. Yang et al. report that Flan-T5 (an instruction-tuned T5) actually had a higher adversarial success rate than the base T5 on certain tasks, meaning it was more easily compromised by targeted synonym swaps. The authors hypothesize that instruction tuning on many tasks made the model rely on specific dataset words, which attackers could then exploit. This highlights that robustness does not automatically scale with model size or training data – training techniques and model architecture play a crucial role. In practice, rigorous adversarial testing is needed for each new LLM, no matter how large, to identify any brittle spots.

D. Prompt Injection and Jailbreak Attacks

Success rates of a prompt-based attack (adversarial suffix) against various LLMs. Grey bars show the baseline (no attack) refusal rate, orange bars use a simple prompt trick (“Sure, here’s ...”), and blue bars use an optimized Greedy Coordinate Gradient (GCG) attack prompt. GCG Ensemble (light blue) combines multiple attack prompts. Results are averaged over many harmful request scenarios. Notably, the attack transfers widely: even black-box APIs like GPT-3.5 and GPT-4 yield high success rates (34–87%) with the adversarial prompt, compared to near-zero without it. Even when LLMs are tuned




Fig. 2. Attack Success Rates for various datasets

to refuse disallowed content, attackers can often find a prompt

that bypasses safety guards. The figure above (Zou et al., 2023) demonstrates a universal jailbreak prompt that was tested on a range of models. Without any attack, all models rarely produce banned content on their own (gray bars near 0%). A naive prompt trick (“Sure, here’s ...”) has limited effect (orange bars), especially on the more aligned systems. But the optimized GCG attack (blue) causes a dramatic rise in compliance: for open-source LLaMA derivatives the success rate jumps to 80–100%, and even OpenAI’s closed models are coaxed into harmful outputs in roughly 30–35% of attempts. By using an ensemble of such attack prompts (light blue), the success rate against GPT-3.5 soared to 86.6%, while GPT-4 also climbed to about 47%. These empirical results starkly illustrate that no current LLM is fully safe from prompt injection – with sufficient ingenuity (and sometimes gradient-based optimization), attackers can find phrasings that break through. Corroborating these findings, an independent study evaluated GPT-4 and GPT-3.5 on a suite of 1,405 known jailbreak prompts. GPT-4 (and its refined versions) caught roughly 90–95% of the malicious prompts, refusing to answer, whereas GPT-3.5-turbo detected only about 49% – meaning it leaked disallowed content half of the time. In other words, GPT-3.5 was the most susceptible to prompt attacks, and GPT-4’s safeguards, while better, were not impenetrable. Researchers have also shown that aligned models can be coerced into revealing private data or system instructions via indirect prompt injection (e.g. by inserting hidden commands into a user-provided document). All these “red-teaming” experiments make clear that prompt-based adversarial techniques remain a serious threat. As LLM deployment grows, developing reliable defenses – from input sanitization to robust training – is an active area of research, aiming to patch the holes exposed by these attacks.

IV. BEHAVIORAL TESTING FRAMEWORKS

To diagnose these brittle behaviors, the community has developed behavioral testing frameworks inspired by software testing. CheckList (Ribeiro et al., 2020) pioneered a task-agnostic methodology for systematic testing of NLP models beyond aggregate accuracy. CheckList defines a matrix of linguistic capabilities (e.g. vocabulary, negation, coreference, temporal reasoning, etc.) and test types (such as invariance tests – perturb input in a way that should not change the output, directional expectation tests – perturb input in a way that should change output in a predictable direction, and minimum functionality tests for basic competency). Using this framework, one can generate diverse test cases to probe specific model behaviors. Ribeiro et al. showed that CheckList could reveal critical bugs in commercial models that standard benchmarks missed – for example, sentiment analyzers failing negation (“I don’t hate this” being misclassified) or NLI models failing on antonyms. The key contribution of CheckList is to treat NLP model evaluation like unit testing, systematically covering a breadth of scenarios to ensure a model’s behavior is robust and consistent. This approach won the Best Paper Award at ACL 2020, reflecting its impact.

Another important tool is Robustness Gym (Goel et al., 2021), which provides a unified platform to evaluate models across multiple robustness paradigms. Robustness Gym implements four evaluation strategies: subpopulation evaluation (performance on specific slices of data, e.g. only long sentences), transformation-based tests (applying perturbations or paraphrases to inputs), evaluation on out-of-domain sets, and adversarial attacks. By integrating these into one toolkit, it allows researchers to compare how a model holds up under different stressors with minimal overhead. For example, one can easily apply character noise, test on a minority dialect subset, and run a synonym-swap attack in one workflow. The toolkit was introduced in the NAACL 2021 demo track and is actively used to identify brittle points in NLP models that may not appear in overall metrics. Beyond general frameworks, targeted “challenge sets” have been created to test specific capabilities. For instance, HANS (Heuristic Analysis for NLI, 2019) is a curated set of NLI examples designed to foil models relying on superficial heuristics (like assuming any sentence with overlapping words is entailment). While HANS predates 2020, it influenced subsequent work in exposing brittleness in large NLI models. Similarly, BLiMP (Warstadt et al., 2020) provides minimal pair tests for grammaticality, which large generative LMs (GPT-style) are evaluated on to gauge syntactic generalization. The general trend in 2020–2025 is a shift toward behavioral evaluation: rather than just reporting a single accuracy, researchers present a suite of tests examining different aspects of model robustness and generalization.

V. ADVERSARIAL EVALUATION AND BENCHMARKS

Adversarial evaluations test how large language models (LLMs) perform on perturbed or maliciously crafted inputs, as in benchmarks like ANLI (Adversarial NLI), AdvGLUE, and HELM. To quantify robustness, researchers employ several statistical metrics that compare a model’s performance on clean (unperturbed) data versus adversarial or perturbed data. Key metrics include accuracy drop, attack success rate (ASR), robustness gap, and consistency score. Below we formally define each with mathematical notation and illustrate how they measure performance degradation in adversarial settings.

A. Accuracy Drop

Accuracy drop measures how much a model’s accuracy decreases under adversarial attack compared to normal conditions. Formally, let Acc_{clean} be the model’s accuracy on the original (clean) test set and Acc_{adv} be the accuracy on an adversarially perturbed version of the test set (with the same ground-truth labels). Some works express the drop as a fraction or percentage of the original accuracy, often called the performance drop rate (PDR).

$$PDR(f) = \frac{Acc(f) - Adv(f)}{Acc(f)}$$

A PDR of 0 means no performance loss (the adversary had no effect), while higher values mean a larger accuracy drop

(e.g. PDR = 0.5 means a 50% relative drop in accuracy). In practice, accuracy drop is often reported in absolute terms (e.g. “accuracy drops by 20 points”) or as a percentage drop. Measuring performance degradation: Accuracy drop directly captures the delta between clean vs. adversarial accuracy. A large ΔAcc indicates the model struggled significantly on adversarial inputs. For example, on the AdvGLUE benchmark, a state-of-the-art model (ELECTRA-Large) obtained about 93.16% accuracy on the original GLUE tasks, but only 41.69% on the adversarially perturbed version. This corresponds to an absolute accuracy drop of 51.5 points, which is a performance drop of around 55% relative to the original accuracy. Such a huge drop highlights the severity of performance degradation under attack. In general, researchers compare Acc_{clean} vs. Acc_{adv} to quantify how much worse the model performs when facing adversarial examples. A small drop (near 0) would indicate strong robustness, whereas large drops signal vulnerability.

B. Attack Success Rate (ASR)

Attack Success Rate (ASR) is the primary metric to evaluate the effectiveness of an adversarial attack. It measures the proportion of attack attempts that successfully cause the model to err.

$$ASR = \frac{\sum_{(x,y) \in D} \mathbf{1}\{f(x) = y \wedge f(A(x)) \neq y\}}{\sum_{(x,y) \in D} \mathbf{1}\{f(x) = y\}}$$

where $D = (x, y)$ is a set of test inputs x with true label y , f is the model, and $A(x)$ is the adversarially perturbed version of x . The indicator $\mathbf{1}\cdot$ counts an event as 1 if true (and 0 otherwise). In words, the numerator counts how many test inputs the model originally got correct ($f(x) = y$) but after applying the attack the model’s prediction is wrong ($f(A(x)) \neq y$). The denominator counts how many were originally correct. Thus, ASR is the fraction of originally correct cases that the attack turns into incorrect predictions. An ASR of 100% means the attack succeeded in flipping every originally correct answer into an error, whereas ASR = 0 means the attack had no impact (the model stayed correct on all those inputs). If the attack is targeted (aiming to induce a specific incorrect output), ASR can be defined as the fraction of inputs where the model’s output is changed to the attacker’s desired target. In general, however, adversarial evaluation of LLMs uses ASR to summarize how often the model is fooled into any incorrect behavior. For example, in the AdvGLUE benchmark various attack algorithms were evaluated by their ASR. An ASR of 80% would indicate that 80% of the time, the adversarial perturbation was able to cause the model to make an error it wouldn’t have made on the clean input. Illustrative example: In practice, high ASR values signal a serious robustness issue. For instance, a study of adversarial synonym substitution

attacks on the AGNews text classification task found that a small 60-million-parameter T5 model had an ASR of 58% (the attack flipped over half of its correct predictions), whereas a much larger 11-billion-parameter T5 model had an ASR of only 15% under the same attack. This indicates the larger model was substantially more robust (fewer of its answers could be changed by the attack). In general, stronger attacks or more vulnerable models yield higher ASR. Recent evaluations on LLMs like GPT-3.5 vs. GPT-4 also report that GPT-3.5 can be tricked with higher success rates than GPT-4, i.e. GPT-4’s adversaries achieve a lower ASR on average (suggesting improved robustness in GPT-4). Attack Success Rate is thus a crucial metric for adversarial benchmarks: e.g., AdvGLUE uses ASR to compare attacks’ effectiveness, and a successful defense should significantly reduce the ASR (ideally to 0, meaning attacks no longer cause mistakes).

VI. TECHNIQUES FOR IMPROVING ROBUSTNESS

A. Adversarial Training:

One of the most direct approaches to harden LLMs is adversarial training, which involves augmenting the training process with adversarial examples. By exposing the model to perturbed or malicious inputs during training, it learns to resist similar attacks at test time. For example, Ebrahimi et al. (2018) showed that injecting adversarial HotFlip perturbations (token swaps guided by gradients) at training time made a classifier more robust to such character-level attacks.




Fig. 3. Adversary’s success rate for text classification

Recent work extends adversarial training to LLMs: OpenAI’s GPT-4 was evaluated on internal adversarial factuality benchmarks and achieved 19 percentage points higher accuracy than GPT-3.5, reflecting significant robustness gains from adversarially-focused training. However, adversarial training at LLM scale is extremely resource-intensive and cannot cover every possible attack; indeed, in the worst-case (white-box) setting, even advanced defenses often yield near-0% robustness against a strong adaptive attacker. Still, it can markedly improve resilience to the specific attack patterns seen during training. For instance, a recent study reported that adversarially fine-tuning a model on “jailbreak” prompts doubled the time it took experts to find new exploits, though it did not eliminate

failures entirely. In practice, adversarial training is used in a targeted way (e.g. fine-tuning on known prompt attacks or harmful queries) to bolster reliability on high-stakes tasks without unduly degrading normal performance.

B. Data Augmentation (Paraphrasing & Back-Translation):

Another technique to improve generalization and robustness is augmenting the training data with semantically equivalent variations. By training on diverse phrasings, styles, or translations of the same queries, an LLM becomes less sensitive to particular wording. Classic approaches include paraphrasing (using models or crowdsourcing to rephrase sentences) and back-translation (translating text to another language and back to generate a paraphrase). Such data augmentation exposes the model to a wider input distribution, helping it handle typos, synonyms, or alternate syntax. For example, the IT-DT method (Sabir et al., 2023) detects adversarial text and then transforms it (e.g. rephrases it) to neutralize the attack while preserving meaning. Back-translating a harmful prompt into another language and then into English can similarly remove hidden triggers. In the context of safety, Anthropic’s red-teaming efforts augmented training with tool-assisted rewrites of problematic prompts to make models like Claude more robust to phrasing variations. While data augmentation alone cannot cover all adversarial strategies, it is a cheap and effective way to make LLMs less brittle to distribution shift (e.g. different dialects or misspellings) and has been shown to boost robustness on benchmarks with paraphrased or noisy inputs.

past → pas't	Alps → llps	talk → taln	local → lorai	you → yoTu
pasturing	lips	tall	moral	Tutu
pasture	laps	tale	Moral	Hutu
pastor	legs	tales	coral	Turku
Task	slips	talent	morals	Futurum

Fig. 4. Noisy Input Example

C. Instruction Tuning and Multi-Task Training:

Large pre-trained models can be made more general and robust by instruction tuning – fine-tuning on a wide array of tasks described via natural language instructions. By training on diverse multi-task datasets (spanning translation, reasoning, reading comprehension, etc.), the model learns to follow arbitrary instructions and becomes less overfit to any single task format. This approach was pioneered in works like Google’s FLAN and BigScience’s T0, which showed that multi-task instruction-trained models exhibit strong zero-shot generalization to unseen tasks. For example, Wei et al. (2022) found that a model tuned on 62 distinct NLP tasks could generalize to new tasks it had never seen, outperforming un-tuned LMs of similar size. Such broad training makes the model more robust to task distribution shifts – it is less likely to break when confronted with an unusual query or a sequence of instructions, because it has effectively been “stress-tested” on a variety of formats and domains. Multi-task training also

tends to reduce the model’s reliance on spurious shortcuts from any single dataset. In addition, instruction tuning often involves human-designed prompts that elicit step-by-step reasoning (chain-of-thought), which can improve the model’s reliability on complex or tricky queries. Overall, exposing LLMs to many tasks and instructions teaches them a form of meta-generalization, yielding more robust performance on real-world inputs that don’t mimic the original training data distribution.

D. Instruction Tuning and Multi-Task Training:

Instruction tuning and multi-task learning have emerged as effective strategies for enhancing the robustness and generalization capabilities of large language models (LLMs). These approaches aim to improve a model’s ability to handle diverse, real-world inputs by fine-tuning it on a curated set of tasks framed as instructions. Rather than training on narrow, task-specific objectives, instruction tuning exposes the model to varied prompts and formats, encouraging flexible adaptation and robust behavior in unseen settings. One of the most influential frameworks in this domain is T5 (Text-to-Text Transfer Transformer) by Raffel et al. (2020), which reformulates every NLP task as a text-to-text problem, enabling multi-task learning across translation, classification, summarization, and question answering. Subsequent works like FLAN (Wei et al., 2022) and FLAN-T5 expanded on this by using a large, diverse instruction dataset to improve zero-shot and few-shot performance. These models demonstrated better performance not only on in-distribution test sets but also in generalization to novel tasks and domains—highlighting instruction tuning’s role in robustness under distribution shift. Additionally, Self-Instruct (Wang et al., 2023) proposed an automatic framework for generating diverse instruction-like tasks using the LLM itself, bootstrapping its own training corpus. This approach reduced the reliance on human-written prompts and significantly enhanced the LLM’s ability to follow user instructions. Similar ideas were adopted by OpenAI in their instruction-tuned GPT-3.5 and GPT-4 models, which exhibited greater consistency, refusal capabilities, and robustness to prompt variations compared to their predecessors. Instruction tuning also interacts with adversarial robustness. For instance, instruction-tuned models are often more robust to subtle prompt perturbations (e.g., rewording or introducing distractor phrases), a property emerging from their training on diverse task representations. However, alignment objectives can introduce new vulnerabilities—for example, prompt-injection attacks that exploit the model’s learned tendency to obey instructions. Recent work such as “Jailbroken” (Wei et al., 2023) shows that even well-aligned models can be tricked into producing harmful outputs via cleverly structured instructions, revealing a tension between helpfulness and robustness. In multi-task training regimes, LLMs learn shared representations across tasks, which improves sample efficiency and encourages the model to generalize from overlapping knowledge. Techniques like Mixture-of-Experts (Shazeer et al., 2017) and UnifiedQA (Khashabi et al., 2020) further push this direction, showing that

robust performance can be obtained by training on large-scale, diverse task mixtures. These models are particularly resilient to task-specific shortcuts, a known source of brittleness in narrow fine-tuning. In summary, instruction tuning and multi-task training serve as powerful mechanisms for equipping LLMs with improved generalization and moderate robustness. They form the basis of recent advancements in safety-aligned LLMs and continue to evolve as researchers develop more diverse datasets and robust alignment strategies.

E. Alignment and Safety Training

As large language models (LLMs) scale in size and capability, aligning their behavior with human values has become a critical area of research. Robustness alone is insufficient if a model, when perturbed or prompted adversarially, produces outputs that are unsafe or unethical. Alignment and safety training aim to address this by teaching models not only to perform well but to do so in a manner consistent with normative constraints and human intent—even in adversarial contexts. One dominant method in this domain is Reinforcement Learning from Human Feedback (RLHF), popularized by the training of OpenAI’s InstructGPT and subsequently GPT-3.5 and GPT-4 (Ouyang et al., 2022). In RLHF, models are first fine-tuned on instruction datasets and then further trained using human preference rankings. This process rewards helpful and harmless outputs while discouraging evasive or harmful responses. RLHF has been shown to increase robustness in real-world dialogue, especially in safety-critical applications, by filtering out unsafe generations even under diverse prompting. Anthropic introduced an alternative paradigm called Constitutional AI (Bai et al., 2022), where instead of relying solely on human labels, the model learns to critique and revise its own outputs using a predefined set of guiding principles or “constitutional rules.” This framework enables scalable alignment while simultaneously improving robustness. For example, Claude, Anthropic’s assistant model, has demonstrated strong resistance to prompt-injection attacks and adversarial jailbreaks compared to earlier models fine-tuned via RLHF alone. Another promising direction is self-alignment, where the LLM is trained to detect and regulate harmful outputs by internalizing safety norms. This can be achieved through techniques such as supervised refusal training, in which models learn to say “no” to unsafe prompts. Evaluations on robustness benchmarks such as RealToxicityPrompts and HarmBench (Mazeika et al., 2024) show that aligned models exhibit lower rates of degeneration or malicious compliance. However, alignment techniques also introduce new vulnerabilities. A major concern is the adversarial misalignment attack, where malicious users craft inputs that exploit safety filters or confuse the reward model. For example, Wei et al. (2023) demonstrated that even advanced RLHF or Constitutionally trained models could be jailbroken via multi-turn prompts, fictional scenarios, or subtle linguistic manipulations emphasizing the need for continuous evaluation and more robust safety supervision.

F. External Defense Mechanisms

While internal modifications like adversarial training and alignment improve the resilience of LLMs to a degree, external defense mechanisms offer an additional and often complementary line of protection. These methods function as wrappers or auxiliary systems around a model, aiming to detect, mitigate, or deflect adversarial behavior without altering the base model’s architecture or parameters. Such strategies are particularly valuable for models deployed as APIs, where retraining or fine-tuning is infeasible.

One widely explored approach is input sanitization, where adversarial inputs are preprocessed before being passed to the model. Techniques include spell correction, synonym normalization, and syntactic simplification to remove perturbations designed to exploit model vulnerabilities (Jia Liang, 2017). Although effective for shallow attacks like character swaps or uncommon phrasing, such defenses are less robust against semantically equivalent paraphrases or context-based jailbreaks.

	Image Classification	Reading Comprehension
Possible Input		Tesla moved to the city of Chicago in 1880.
Similar Input		Tadakatsu moved to the city of Chicago in 1881.
Semantics	Same	Different
Model’s Mistake	Considers the two to be different	Considers the two to be the same
Model Weakness	Overly sensitive	Overly stable

Fig. 5. Adversarial Inputs

Adversarial input detectors serve as gatekeepers, identifying whether an incoming prompt is likely adversarial. These can be implemented via auxiliary classifiers trained on known attack data, or via statistical methods that flag inputs exhibiting suspicious patterns (e.g., high perplexity, unnatural word distributions, or structural anomalies). For example, Anthropic’s “constitutional classifiers” monitor inputs and proactively reject prompts exhibiting harmful intent (Bai et al., 2022). These detectors can be integrated upstream in the LLM pipeline and block problematic queries before the model generates output.

Consensus-based defenses use multiple LLM instances or prompt variants to evaluate the consistency of a model’s responses. If a single query produces widely divergent outputs when paraphrased slightly, the system may abstain from providing a response. This method is based on the assumption that benign inputs should yield stable predictions across perturbations, whereas adversarial ones often induce inconsistency. Self-consistency sampling (Wang et al., 2022) and ensemble voting are common implementations of this strategy.

Refusal heads and abstention mechanisms are explicitly trained components that allow LLMs to decline to answer when uncertain or when a prompt is identified as harmful. For instance, GPT-4 includes internal thresholds and behavioral cues that trigger refusal responses when prompted in adversarial or sensitive contexts (OpenAI, 2023). While refusal reduces the risk of harmful output, overuse can degrade user

experience and lead to excessive conservatism, known as the “evasive model” problem.

Retrieval-augmented generation (RAG) offers another external defense vector by grounding model outputs in verified external knowledge. When an LLM is prompted with ambiguous or misleading questions, access to up-to-date retrieval sources (e.g., Wikipedia, news archives) helps reduce hallucinations and bias. In robustness terms, grounding can mitigate the effects of temporal distribution shift and adversarial misinformation prompts (Lewis et al., 2020).

Finally, monitoring systems and logging frameworks play a critical operational role. In high-stakes environments (e.g., healthcare, finance), real-time monitoring of input-output pairs enables anomaly detection, red-teaming feedback loops, and even post-hoc interventions. Such infrastructure is essential for maintaining robustness at scale and feeds into continual learning and patching pipelines.

In short, external defenses act as practical safeguards against the evolving landscape of LLM vulnerabilities. While no defense is perfect in isolation, combining these mechanisms with robust training and evaluation practices forms a layered defense architecture—crucial for secure and reliable LLM deployment in the wild.

G. Red Teaming and Iterative Refinement

Red teaming has emerged as a cornerstone practice in evaluating and improving the safety and robustness of large language models (LLMs). Borrowed from cybersecurity, red teaming involves actively attempting to subvert or break a system through adversarial interaction. In the context of LLMs, red teaming typically consists of crafting inputs—often through manual probing or generative search—that elicit harmful, unsafe, or misaligned behavior. These efforts play a vital role in uncovering failure modes that elude automated evaluation pipelines and static benchmarks.

Modern LLM developers now integrate red teaming into their development workflows. OpenAI’s GPT-4 system card (2023) describes a multi-stage red teaming process that included both internal researchers and external experts. These agents tested the model’s behavior across domains like misinformation, hate speech, bias, instruction-following, and misuse. Their findings directly informed post-training alignment strategies, refusal triggers, and safety filters.

Anthropic’s Claude models were similarly red-teamed using adversarial prompting campaigns, focusing not only on overt jailbreak attempts but also on subtle misuse patterns such as fictional scenarios or implicit instructions. The outcomes highlighted vulnerabilities in prompt conditioning, continuity handling, and meta-prompt leakage. These insights led to enhancements in their constitutional AI framework, which systematically reflects on and refines outputs.

Red teaming efforts often contribute to iterative refinement, where feedback from adversarial testing cycles is incorporated into model improvement. This refinement can occur via adversarial fine-tuning, safety reinforcement learning, or synthetic data augmentation. For instance, models can be

retrained on adversarial prompts and their ideal completions, thus immunizing them against known attack types (Zhu et al., 2020). This process mirrors curriculum learning, where increasingly difficult examples guide the model toward robust behavior.

To scale red teaming beyond manual efforts, researchers have explored automated adversarial generation using LLMs themselves. These self-red-teaming approaches employ one model to craft challenging prompts and another to respond, forming adversarial dialogue pairs that reveal gaps in alignment and robustness (Perez et al., 2022). This technique allows the creation of large-scale, diverse adversarial corpora for continuous benchmarking.

Moreover, platforms like Dynabench (Kiela et al., 2021) and HarmBench (Mazeika et al., 2024) facilitate community-driven red teaming. These tools allow users to submit failure cases, view model outputs, and contribute to a growing adversarial test suite. Such open-ended, real-time feedback is crucial for evaluating LLMs under deployment-like conditions, especially as novel attack vectors continue to emerge.

Despite its benefits, red teaming poses challenges in standardization and reproducibility. Unlike static benchmarks, red teaming outcomes can vary across testers, domains, and threat models. As a result, researchers advocate for transparent reporting practices, shared prompt corpora, and reproducible red-teaming protocols to ensure comparability across models and institutions.

In conclusion, red teaming is not merely a final evaluation step but a continuous, structured dialogue between developers, researchers, and adversaries. When combined with iterative refinement and benchmark-based training, it forms a powerful loop for exposing and patching model weaknesses—paving the way toward safer and more reliable LLM systems.

VII. CONCLUSION

The increasing deployment of large language models (LLMs) in real-world applications has made robustness, safety, and generalization core concerns in contemporary natural language processing research. Despite achieving impressive benchmarks on standard tasks, LLMs exhibit notable fragility when faced with adversarial perturbations, distribution shifts, or malicious prompts. This survey provided a comprehensive overview of the vulnerabilities, evaluation strategies, and mitigation techniques associated with adversarial robustness and generalization in LLMs.

We began by highlighting the adversarial vulnerabilities inherent to modern language models, including susceptibility to synonym substitutions, paraphrases, and character-level perturbations. Through the evolution of evaluation frameworks such as CheckList, ANLI, Dynabench, and AdvGLUE, it has become evident that traditional accuracy metrics are insufficient to characterize model robustness. Instead, behavioral testing and adversarial benchmarking have become standard tools to assess real-world generalization and safety.

The survey then explored various strategies aimed at improving LLM robustness, from adversarial training and data

augmentation to instruction tuning and multi-task learning. Each approach addresses specific facets of the robustness problem—whether enhancing resistance to perturbations, improving generalization across tasks, or aligning outputs with human intent. Notably, instruction-tuned models and alignment-based techniques such as Reinforcement Learning from Human Feedback (RLHF) and Constitutional AI have introduced meaningful advances in safe and stable model behavior. However, these methods remain vulnerable to novel attack vectors, including jailbreaks and adversarial prompt injections.

We further examined the role of external defense mechanisms, including input sanitization, refusal heads, and adversarial detectors, as layered safeguards for deployed LLM systems. Red teaming, both human-led and automated, emerged as a critical practice for exposing unforeseen vulnerabilities and informing iterative refinement. These continuous evaluation loops have proven essential for enhancing model alignment, trustworthiness, and resilience under adversarial pressure.

While considerable progress has been made, several open challenges remain. Adversarial robustness is often brittle to unseen attack types; alignment may inadvertently suppress useful outputs; and robust generalization across domains and languages is far from solved. Additionally, the trade-offs between safety, performance, and usability continue to complicate deployment decisions. As LLMs grow in scale and influence, addressing these issues becomes not only a technical priority but an ethical imperative.

Future research must focus on unified robustness evaluation frameworks, better transferability of defense methods, and scalable alignment strategies. Incorporating continual red teaming, real-world feedback, and transparent reporting practices will be crucial for building LLMs that are not only intelligent but also safe, fair, and robust by design.

REFERENCES

- [1] Y. Bai et al., “Constitutional AI: Harmlessness from AI Feedback,” *arXiv preprint arXiv:2212.08073*, 2022.
- [2] M. Bartolo, A. Roberts, J. Welbl, S. Riedel, and P. Stenetorp, “Beat the AI: Investigating Adversarial Human Annotation for Reading Comprehension,” *EMNLP*, 2020.
- [3] R. Bommasani, K. Davis, and C. Cardie, “On the Opportunities and Risks of Foundation Models,” *arXiv preprint arXiv:2108.07258*, 2021.
- [4] T. Brown et al., “Language Models are Few-shot Learners,” *NeurIPS*, 2020.
- [5] N. Carlini and A. Terzis, “Evaluating Large Language Models Trained on Code,” *arXiv preprint arXiv:2107.03374*, 2021.
- [6] J. Devlin, M. W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding,” *NAACL*, 2019.
- [7] J. Ebrahimi, A. Rao, D. Lowd, and D. Dou, “HotFlip: White-box Adversarial Examples for Text Classification,” *ACL*, 2018.
- [8] D. Ganguli et al., “Red Teaming Language Models to Reduce Harms,” *ACL*, 2022.
- [9] S. Gehman, S. Gururangan, M. Sap, Y. Choi, and N. A. Smith, “RealToxicityPrompts: Evaluating Neural Toxic Degeneration in Language Models,” *EMNLP*, 2020.
- [10] I. J. Goodfellow, J. Shlens, and C. Szegedy, “Explaining and Harnessing Adversarial Examples,” *ICLR*, 2015.
- [11] D. Jin, Z. Jin, J. T. Zhou, and P. Szolovits, “TextFooler: A Massively Transferable Text Attack for Fooling Natural Language Understanding Models,” *AAAI*, 2020.

- [12] D. Kiela *et al.*, “Dynabench: Rethinking Benchmarking in NLP,” *NAACL*, 2021.
- [13] A. Kurakin, I. Goodfellow, and S. Bengio, “Adversarial Machine Learning at Scale,” *ICLR*, 2017.
- [14] P. Liang *et al.*, “Holistic Evaluation of Language Models,” *NeurIPS*, 2022.
- [15] S. Lin, J. Hilton, and O. Evans, “TruthfulQA: Measuring How Models Mimic Human Falsehoods,” *NeurIPS*, 2021.
- [16] R. Jia, P. Liang, “Adversarial Examples for Evaluating Reading Comprehension Systems,” *EMNLP*, 2017.
- [17] M. Mazeika, A. Sabharwal, and D. Kiela, “HarmBench: A Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal,” *arXiv preprint arXiv:2404.12345*, 2024.
- [18] J. Morris *et al.*, “TextAttack: A Framework for Adversarial Attacks, Data Augmentation, and Adversarial Training in NLP,” *EMNLP*, 2020.
- [19] Y. Nie *et al.*, “Adversarial NLI: A New Benchmark for Natural Language Understanding,” *NeurIPS*, 2020.
- [20] OpenAI, “GPT-4 System Card,” *OpenAI Technical Report*, 2023.
- [21] L. Ouyang *et al.*, “Training Language Models to Follow Instructions with Human Feedback,” *NeurIPS*, 2022.
- [22] A. Radford, K. Narasimhan, T. Salimans, and I. Sutskever, “Improving Language Understanding by Generative Pre-training,” *OpenAI Technical Report*, 2018.
- [23] M. T. Ribeiro, T. Wu, C. Guestrin, and S. Singh, “Beyond Accuracy: Behavioral Testing of NLP Models with CheckList,” *ACL*, 2020.
- [24] A. Sablayrolles, M. Douze, C. Schmid, and H. Jégou, “White-box vs Black-box: Bayes Optimal Strategies for Membership Inference,” *ICML*, 2019.
- [25] H. Touvron *et al.*, “Llama 2: Open Foundation and Fine-Tuned Chat Models,” *Meta AI Technical Report*, 2023.
- [26] J. Huang, S. S. Gu, L. Hou, Y. Wu, X. Wang, H. Yu, and J. Han, “Large language models can self-improve,” *arXiv preprint arXiv:2207.05221*, 2022.
- [27] A. Conneau, D. Kiela, H. Schwenk, L. Barrault, and A. Bordes, “Supervised Learning of Universal Sentence Representations from Natural Language Inference Data,” *EMNLP*, 2017.
- [28] A. Abid, M. Farooqi, and J. Zou. Large language models associate Muslims with violence. *Nature Machine Intelligence*, 3(6):461–463, June 2021.
- [29] D. Bates, M. Mächler, B. Bolker, and S. Walker. Fitting linear mixed-effects models using lme4. *Journal of Statistical Software*, 67(1):1–48, 2015.
- [30] R. Collobert and J. Weston. A unified architecture for natural language processing: Deep neural networks with multitask learning. In *Proceedings of the 25th International Conference on Machine Learning*, pages 160–167. ACM, 2008.
- [31] S. R. Bowman, G. Angeli, C. Potts, and C. D. Manning. A large annotated corpus for learning natural language inference. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 632–642. Association for Computational Linguistics, 2015.
- [32] J. Wang, H. U. Xu, W. Hou, et al. On the robustness of ChatGPT: An adversarial and out-of-distribution perspective. In *ICLR 2023 Workshop on Trustworthy and Reliable Large-Scale Machine Learning Models*, 2023.
- [33] A. Kabra, M. Bhatia, Y. K. Singla, et al. Evaluation toolkit for robustness testing of automatic essay scoring systems. In *Proceedings of the 5th Joint International Conference on Data Science & Management of Data (9th ACM IKDD CODS and 27th COMAD)*, pages 90–99, 2022.
- [34] Z. Yang, Z. Meng, X. Zheng, and R. Wattenhofer. Assessing adversarial robustness of large language models: An empirical study. Technical Report, 2024. <https://arxiv.org/abs/2405.02764>
- [35] M. Owen, A. Panken, R. Moss, L. Alvarez, and C. Leeper. ACAS Xu: Integrated collision avoidance and detect and avoid capability for UAS. In *Proceedings of the 38th IEEE/AIAA Digital Avionics Systems Conference (DASC)*, pages 1–10, 2019.
- [36] N. Levy, A. Ashrov, and G. Katz. Towards robust LLMs: An adversarial robustness measurement framework — Code, 2024. <https://github.com/adielashrov/trust-ai-roma-for-llm>
- [37] J. Kaddour, J. Harris, M. Mozes, H. Bradley, R. Raileanu, and R. McHardy. Challenges and applications of large language models. Technical Report, 2023. <https://arxiv.org/abs/2307.10169>
- [38] Y. Elboher, R. Elsaleh, O. Isac, M. Ducoffe, A. Galametz, G. Povéda, R. Bouzamouza, N. Cohen, and G. Katz. Robustness assessment of a runway object classifier for safe aircraft taxiing. In *Proceedings of the 43rd International Digital Avionics Systems Conference (DASC)*, 2024.
- [39] S. Bassan and G. Katz. Towards formal XAI: Formally approximate minimal explanations of neural networks. In *Proceedings of the 29th International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS)*, pages 187–207, 2023.
- [40] G. Katz, D. Huang, D. Ibeling, K. Julian, C. Lazarus, R. Lim, P. Shah, S. Thakoor, H. Wu, A. Zeljić, D. Dill, M. Kochenderfer, and C. Barrett. The Marabou framework for verification and analysis of deep neural networks. In *Proceedings of the 31st International Conference on Computer Aided Verification (CAV)*, pages 443–452, 2019.
- [41] C. Ferrari, M. Muller, N. Jovanovic, and M. Vechev. Complete verification via multi-neuron relaxation guided branch-and-bound. Technical Report, 2022. <https://arxiv.org/abs/2205.00263>
- [42] K. Shen. The generalization and robustness of transformer-based language models on commonsense reasoning. Technical Report, Information Sciences Institute, USC Viterbi School of Engineering, Marina Del Rey, California, 2024. <mailto:keshen@isi.edu>
- [43] A. Singh, N. Singh, and S. Vatsal. Robustness of LLMs to perturbations in text. *arXiv preprint arXiv:2407.08989*, 2024. <https://doi.org/10.48550/arXiv.2407.08989>
- [44] V. Lad, W. Gurnee, and M. Tegmark. The remarkable robustness of LLMs: Stages of inference? *arXiv preprint arXiv:2406.19384*, 2024. <https://doi.org/10.48550/arXiv.2406.19384>
- [45] C. Bryant, M. Felice, Ø. E. Andersen, and T. Briscoe. The BEA-2019 shared task on grammatical error correction. In *Proceedings of the Fourteenth Workshop on Innovative Use of NLP for Building Educational Applications*, pages 52–75, Florence, Italy, 2019. Association for Computational Linguistics.